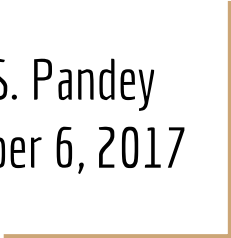




Introduction to Java Concurrency

Anand S. Pandey
September 6, 2017



Agenda

Processes vs. Threads. Why threads ?

Java threads. Thread creation approaches.

What is Thread Safety ? Sharing Objects between threads

Synchronization approaches. Volatile and Final variables

Deadlock examples.

ExecutorService + Future + Callable example

Java Memory Model (JMM). Happens-Before relationship.

Process vs. Threads

- **Process** is a program in execution.
- **Thread** is a path of execution ***within*** a process. It is a basic unit of CPU utilization. It comprises of:
 - thread Id
 - a program counter
 - a register set
 - a stack
- **Thread** shares with other threads belonging to the **same** process, its *code* section, *data* section and other operating system resources.
 - See picture in next slide

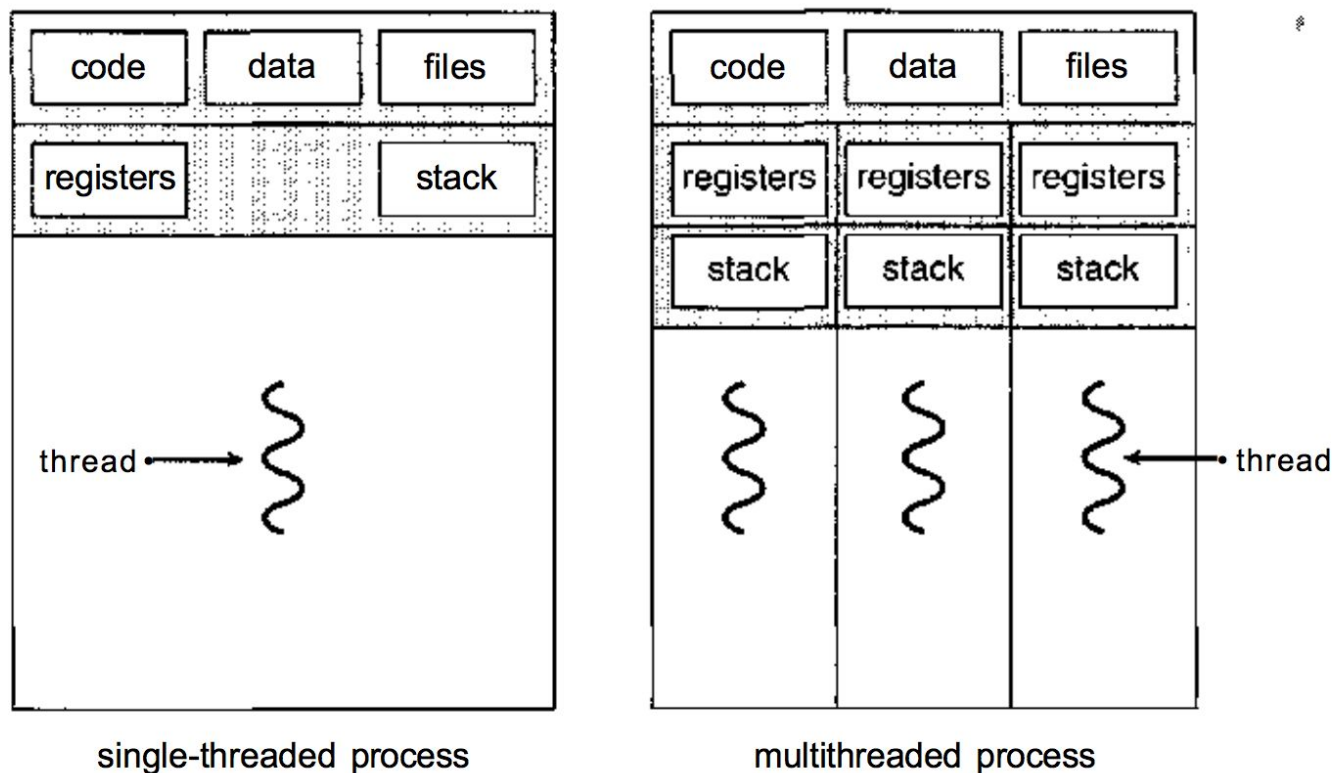


Figure 4.1 Single-threaded and multithreaded processes.

Benefits of Multithreaded programming


- Responsiveness
 - Allows a program to continue running even if parts of it is “waiting”
- Resource sharing
 - Threads share memory and resources of the process to which they belong
 - All threads run within same address space
- Economy
 - They can communicate through shared data and eliminate the overhead of system calls
- Utilization of multiprocessor architectures
 - They allow you to get parallel performance on a shared memory multiprocessor.

Java Threads

- **Threads** are the fundamental model of program execution in a Java program.
- Java language and its APIs provide a rich set of features for the creation and management of threads
- Various approaches of creating threads:
 - By extending **Thread** class and overriding **run()** method
 - Using **Runnable** interface
 - Using **ExecutorService**

Approach 1 (extending Thread class)

```
package main.demo;

class MyThread extends Thread {
    
    public void run() {
        System.out.println("Inside my thread");
    }
}

public class ThreadDemo {

    public static void main(String[] args) {
        MyThread myThread = new MyThread();
        myThread.start();
    }
}
```

Threads – States

When start() method is invoked on thread
It is said to be in Runnable state.
But it is not actually executing the run method.
It is ready to run.

t.start()



Runnable

Sleep/
Waiting/
Blocked

t.sleep(long x)

Running

DEAD

- When the sleep method is invoked, the thread may go to sleep state.
- In sleep, wait or blocked state, it is not running any code. But the thread is alive.
- A thread can move to Wait or blocked state any time.
- A thread in sleep, blocked or wait state can move to Runnable or Running state.
- Invoking sleep method will not guarantee that the thread will go to sleep state immediately

Thread t = new Thread()

When a new thread object is created, the Thread is said to be in new state.

In Running state, the thread is actually executing the code.

The thread has finished executing the run method.
It moves to Dead state.
Note: You cannot use t.start() again on this object once the thread is dead.

Approach 1 continued ...

- Creating a **Thread** object does not specifically create the new thread.
- It is the **start()** method that actually creates the new thread.
- Calling the **start()** method for the new object does two things:
 - It allocates memory and initializes a new thread in the JVM.
 - It calls the **run()** method, making the thread eligible to be run by the JVM.
 - We never call the **run()** method directly. Rather, we call the **start()** method, and it calls the **run()** method on our behalf.

Approach 2 (Using Runnable interface) ...

```
package main.demo;

class MyRunnable implements Runnable {

    @Override
    public void run() {
        for (int i = 0; i < 6; i++) {
            System.out.println("The counter value is " + i);
        }
    }
}
```

```
public class ThreadRunnableDemo {

    public static void main(String[] args) {
        MyRunnable r = new MyRunnable();
        Thread foo = new Thread(r);
        Thread bar = new Thread(r);
        foo.start();
        bar.start();
    }
}
```

Approach 3

Using **ExecutorService**.

Discussed later in this presentation with example code.

Thread Safety

- Writing thread safe programs, at its core, is all about managing access to **state** and in particular to **shared** , **mutable** state.
- An object's **state** is its **data**, stored in state variables viz. instance/static variables
 - Shared → Variable is accessed by multiple threads at a time.
 - Mutable → State of variable can change over period of its lifetime
- We may talk about Thread Safety as if it were about **code**, but it's all about protecting **data** from uncontrolled concurrent access.

- Whenever more than one thread accesses a given state variable, and one of them might **write** to it, they all **must** coordinate their access to it using ***synchronization***.
- The primary mechanism for ***synchronization*** in Java is the ***synchronized*** keyword, which provides *exclusive locking* and *memory visibility*, but the term “synchronization” also ***includes the use of volatile variables, explicit locks, and atomic variables.***

Synchronization and Locks

- Java provides two types of synchronization:
 - Method level synchronization
 - Block level synchronization
- **synchronized** methods prevent more than one thread from accessing an object's critical method code *simultaneously*.
- To synchronize a **block** of code, we must specify an argument that is the object whose lock you want to synchronize on.
- While only *one* thread can be accessing **synchronized** code of a particular instance, multiple threads can still access the same object's **unsynchronized** code.

- When a thread goes to **sleep**, its locks will be *unavailable* to other threads.
- *static* methods can be synchronized, using the lock from the *java.lang.Class* instance representing that class.

Defining Thread Safety

- A class is *thread-safe* when it continues to behave correctly when accessed from multiple threads.
 - Correctness means that a class conforms to its specification.
 - A good specification defines **invariants** constraining an object's state and **postconditions** describing the effects of its operations.
- A class is ***thread-safe*** if it behaves *correctly* when accessed from multiple threads, regardless of the scheduling or interleaving of the execution of those threads by the runtime environment, and ***with no additional synchronization or other coordination on the part of the calling code.***

Synchronizing Code

- Account withdrawal example (IDE)

Sharing Objects

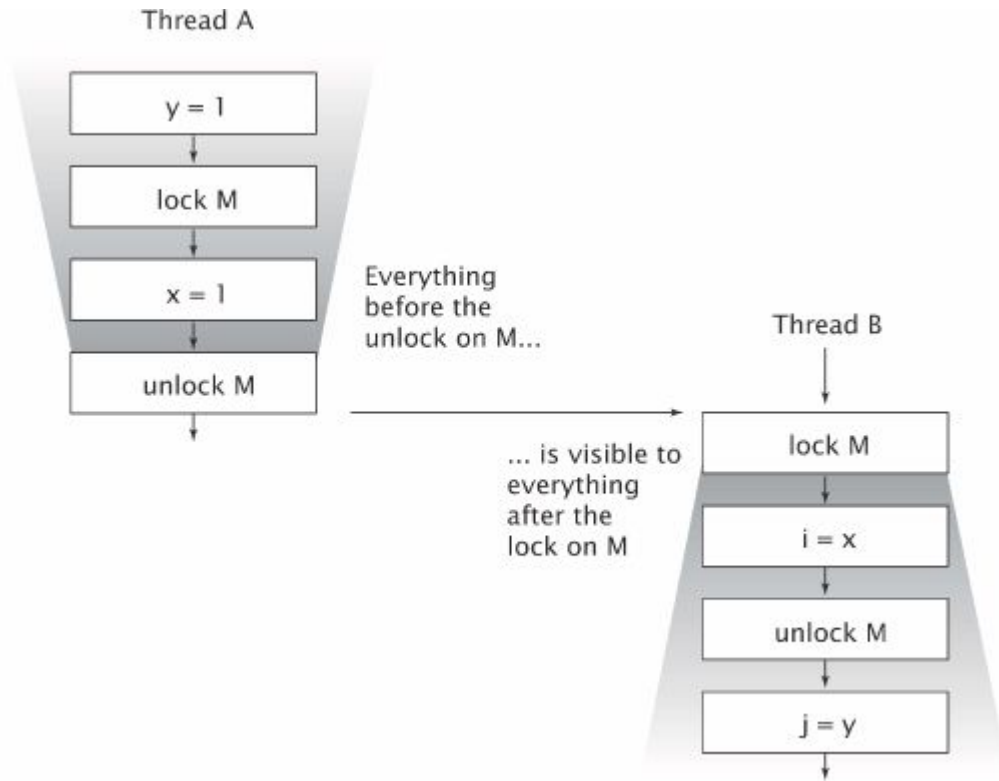
```
public class NoVisibility {  
    private static boolean ready;  
    private static int number;  
  
    private static class ReaderThread extends Thread {  
        public void run() {  
            while (!ready)  
                Thread.yield();  
            System.out.println(number);  
        }  
    }  
  
    public static void main(String[] args) {  
        new ReaderThread().start();  
        number = 42;  
        ready = true;  
    }  
}
```



Some possibilities ...

- **NoVisibility** could loop forever because the value of **ready** might never become visible to the *Reader* thread.
- **NoVisibility** could print **zero** because the write to **ready** might be made visible to the *Reader* thread before the write to **number**, a phenomenon known as **reordering**.
- When the *Main* thread writes first to **number** and then to **ready** **without synchronization**, the *Reader* thread could see those writes happen **in the opposite order -- or not at all**.

- In the absence of **synchronization**, the compiler, processor, and runtime can do some downright weird things to the order in which operations appear to execute.
- Compilers are allowed to *reorder* the instructions in either thread, when this **does not** affect the execution of that thread in *isolation*.
- Attempts to reason about the order in which memory actions “must” happen in insufficiently synchronized multithreaded programs will almost certainly be incorrect.
- **Always use the proper synchronization whenever data is shared across threads.**




- Intrinsic locking can be used to **guarantee** that one thread sees the effects of another in a ***predictable*** manner

- When thread A executes a **synchronized** block, and subsequently thread B enters a **synchronized** block ***guarded by the same lock***, the values of variables that were visible to A prior to releasing the lock are guaranteed to be visible to B upon acquiring the lock.
- In other words, everything A did in or prior to a **synchronized** block is visible to B when it executes a **synchronized** block ***guarded by the same lock***.
- ***Without synchronization, there is no such guarantee.***
- IDE mode (more on @GuardedBy)

Volatile variables

- When a field is declared ***volatile***, the compiler and runtime are put on notice that this variable is shared and *that operations on it should not be reordered with other memory operations*.
- Volatile variables are *not cached* in registers or in caches where they are hidden from other processors.
- Read of a volatile variable always returns the most recent write by any thread.



Thread A:

```
sharedObject.nonVolatile = 123;
```

```
sharedObject.counter = sharedObject.counter + 1;
```

Thread B:

```
int counter = sharedObject.counter;
```

```
int nonVolatile = sharedObject.nonVolatile;
```

- Example code (IDE mode). Counter Class


```
class VolatileExample {  
    int x = 0;  
    volatile boolean v = false;  
    public void writer() {  
        x = 42;  
        v = true;  
    }  
  
    public void reader() {  
        if (v == true) {  
            ??  
        }  
    }  
}
```

Final fields in Java

- Declaring the variable that is initialized to the object as *final* prevents the object from being *partially initialised*.
- An object is considered to be **completely** initialized when its constructor finishes.
 - A thread that can only see a *reference* to an object after that object has been completely initialized is **guaranteed** to see the *correctly initialized* values for that object's *final* fields.
- Variables that are declared *static* and initialized at declaration OR from a static initializer are **guaranteed** to be **fully constructed** before being made visible to other threads. However, this solution forgoes the benefits of lazy initialization.

```
class FinalFieldExample {  
  
    final int x;  
    int y;  
    static FinalFieldExample f;  
    public FinalFieldExample() {  
        x = 3;  
        y = 4;  
    }  
  
    static void writer() {  
        f = new FinalFieldExample();  
    }  
  
    static void reader() {  
        if (f != null) {  
            int i = f.x;  
            int j = f.y;  
        }  
    }  
}
```


Deadlock Example 1

// Warning: deadlock-prone!

```
public class LeftRightDeadlock {  
    private final Object left = new Object();  
    private final Object right = new Object();
```

```
    public void leftRight() {  
        synchronized (left) {  
            synchronized (right) {  
                doSomething();  
            }  
        }  
    }  
}
```

```
    public void rightLeft() {  
        synchronized (right) {  
            synchronized (left) {  
                doSomethingElse();  
            }  
        }  
    }  
}
```



Deadlock Example 2

// Warning: deadlock-prone!

```
public void transferMoney(Account fromAccount,
                          Account toAccount,
                          DollarAmount amount)
    throws InsufficientFundsException {
    synchronized (fromAccount) {
        synchronized (toAccount) {
            if (fromAccount.getBalance().compareTo(amount) < 0)
                throw new InsufficientFundsException();
            else {
                fromAccount.debit(amount);
                toAccount.credit(amount);
            }
        }
    }
}
```



A: transferMoney(myAccount, yourAccount, 10);

B: transferMoney(yourAccount, myAccount, 20);

How to avoid deadlock ?

Impose ordering on lock acquisition (see next slide)

```
public class MoneyTransfer {
```

```
    private static final Object tieLock = new Object();
```

```
    public void transferMoney(final Account fromAcct, final Account toAcct, double amount) {
```

```
        class Helper {
```

```
            private void transfer() throws InsufficientFundsException {
```

```
                if (fromAcct.getBalance().compareTo(amount) < 0) {
```

```
                    throw new InsufficientFundsException("Insufficient funds to perform transfer");
```

```
                } else {
```

```
                    fromAcct.debitAmount(amount);
```

```
                    toAcct.creditAmount(amount);
```

```
                }
```

```
            }
```

```
        }
```

```
        int fromHash = System.identityHashCode(fromAcct);
```

```
        int toHash = System.identityHashCode(toAcct);
```

```
        if (fromHash < toHash) {
```

```
            synchronized (fromAcct) {
```

```
                synchronized (toAcct) {
```

```
                    new Helper().transfer();
```

```
                }
```

```
            }
```

```
        } else if (toHash < fromHash) {
```

```
            synchronized (toAcct) {
```

```
                synchronized (fromAcct) {
```

```
                    new Helper().transfer();
```

```
                }
```

```
            }
```

```
        } else {
```

```
            synchronized (tieLock) {
```

```
                synchronized (fromAcct) {
```

```
                    synchronized (toAcct) {
```

```
                        new Helper().transfer();
```

```
                    }
```

```
                }
```

```
            }
```

```
        }
```

```
    }
```

Limitations of intrinsic locking

- It is **not** possible to *interrupt* a thread waiting to acquire a lock.
- It is **not** possible to acquire a lock without being willing to wait for it forever.
- Intrinsic locks also **must be** released in the same block of code in which they are acquired.
 - this simplifies coding and interacts nicely with exception handling, but makes *non-block structured* locking disciplines impossible.

Solution: Explicit Locks

```
public interface Lock {  
    void lock();  
    void lockInterruptibly() throws InterruptedException;  
    boolean tryLock();  
    boolean tryLock(long timeout, TimeUnit unit)  
        throws InterruptedException;  
    void unlock();  
    Condition newCondition();  
}
```

ReentrantLock implements **Lock** interface providing same **mutual exclusion** and **memory-visibility** guarantees as **synchronized**.

Example Code (IDE mode)

Template for using Lock

```
Lock lock = new ReentrantLock();  
...  
lock.lock();  
try {  
    // update object state  
    // catch exceptions and restore invariants if necessary  
} finally {  
    lock.unlock();  
}
```

Example (IDE mode) (AccountReentrantLock)

Another example of Lock ...

```
public boolean trySendOnSharedLine(String message,  
                                   long timeout, TimeUnit unit)  
    throws InterruptedException {  
    long nanosToLock = unit.toNanos(timeout)  
        - estimatedNanosToSend(message);  
    if (!lock.tryLock(nanosToLock, NANOSECONDS))  
        return false;  
    try {  
        return sendOnSharedLine(message);  
    } finally {  
        lock.unlock();  
    }  
}
```

```

public boolean transferMoney(Account fromAcct,
                             Account toAcct,
                             DollarAmount amount,
                             long timeout,
                             TimeUnit unit)
    throws InsufficientFundsException, InterruptedException {
    long fixedDelay = getFixedDelayComponentNanos(timeout, unit);
    long randMod = getRandomDelayModulusNanos(timeout, unit);
    long stopTime = System.nanoTime() + unit.toNanos(timeout);

    while (true) {
        if (fromAcct.lock.tryLock()) {
            try {
                if (toAcct.lock.tryLock()) {
                    try {
                        if (fromAcct.getBalance().compareTo(amount)
                            < 0)
                            throw new InsufficientFundsException();
                        else {
                            fromAcct.debit(amount);
                            toAcct.credit(amount);
                            return true;
                        }
                    } finally {
                        toAcct.lock.unlock();
                    }
                }
            } finally {
                fromAcct.lock.unlock();
            }
        }
        if (System.nanoTime() > stopTime)
            return false;
        NANOSECONDS.sleep(fixedDelay + rnd.nextLong() % randMod);
    }
}

```

Code example

ExecutorService + Future + Callable (IDE mode)

Java Memory Model (JMM)

- Java Memory Model (JMM) describes what behaviors are legal in multithreaded code, and how threads may interact through memory (RAM).
- JMM is specified in terms of **actions**, which include *reads* and *writes* to variables, *locks* and *unlocks* of monitors and *starting* and *joining* with threads.
- JMM defines the behavior of *volatile* and *synchronized*, and, more importantly, ensures that a correctly synchronized Java program runs **correctly** on **all** processor architectures.

Happens Before relationship

- The JMM defines a partial ordering called ***happens before*** on all actions within the program.
- To guarantee that the thread executing action B can see the results of action A (whether or not A and B occur in different threads), there **must be** a ***happens before*** relationship between A and B.
 - In the absence of a *happens-before* ordering between two operations, the JVM is free to *reorder* them as it pleases.
- It should be noted that the presence of a *happens before* relationship between two actions does not necessarily imply that they have to take place in that order in an implementation.
 - If the reordering produces results consistent with a legal execution, it is not illegal

Examples of *Happens Before* relationship

- **Monitor lock rule.** An *unlock* on a monitor lock ***happens before*** every subsequent *lock* on that same monitor lock.
- **Volatile variable rule.** A *write* to a volatile field ***happens before*** every subsequent *read* of that same field.
- **Thread start rule.** A call to *Thread.start()* on a thread ***happens before*** every action in the started thread.
- **Thread termination rule.** Any action in a thread ***happens before*** any other thread detects that thread has terminated, either by successfully return from *Thread.join()* or by *Thread.isAlive()* returning false.

- **Interruption rule.** A thread calling `interrupt` on another thread ***happens before*** the interrupted thread detects the interrupt (either by having *`InterruptedException`* thrown, or invoking *`isInterrupted`* or *`interrupted`*).
- **Finalizer rule.** The end of a constructor for an object ***happens before*** the start of the finalizer for that object.
- **Transitivity.** If A ***happens before*** B, and B ***happens before*** C, then A ***happens before*** C.

Some topics for future brown bag sessions...

- Concurrent data structures
 - ConcurrentHashMap, Blocking queues, AtomicX classes
 - Synchronizers *viz.* semaphores, countdownLatch, cyclic barriers, exchanger, phaser.
- ThreadPools. Executors vs. ExecutorService. Various implementations.
- Different **Lock** variations
 - Reentrant locks, ReentrantReadWrite locks, Stamped locks (introduced in Java 8)
- Thread interruption mechanisms
- Non-blocking synchronization algorithms
 - Use of CAS (CompareAndSwap) hardware instruction

References

Java Concurrency in Practice by *Joshua Bloch, Doug Lea et.al.*

Operating System Principles, Seventh Edition by *Galvin, Gagne et.al*

<https://www.securecoding.cert.org/confluence/pages/viewpage.action?pageId=18581044>

Java Language Specification <https://docs.oracle.com/javase/specs/jls/se7/html/jls-17.html>

<http://www.cs.umd.edu/~pugh/java/memoryModel/jsr-133-faq.html>

<https://stackoverflow.com/questions/16213443/instruction-reordering-happens-before-relationship-in-java>

<http://tutorials.jenkov.com/java-concurrency/volatile.html>

Questions ?